CardioTriage AI

Machine Learning Phase — Detailed Walkthrough

From raw cardiac records to a deployed risk-and-phenotype engine (Steps A1–A8)

920	**3**	**0.91**	**0.96**
patients	models compared	ROC-AUC	recall (tuned)

BS Artificial Intelligence · Combined Assessment (Intro to AI + Machine Learning)

CONTEXT

Where the Machine Learning Phase Fits

1 · Data & ML	ML risk engine — predicts per-patient risk probability + phenotype
2 · Classical AI core	A* search · CSP allocator · forward-chaining knowledge base
3 · Orchestration	Gemini + LangChain agent — decides which tool to call & explains
4 · Presentation	Interactive web triage board (ranking, paths, allocations)

This deck covers Layer 1

Phase A builds the ML risk engine — Steps A1–A8.

Its outputs (risk probability + phenotype) become the inputs every layer above consumes.

The LLM never makes a medical call alone — it coordinates deterministic tools.

DATA

The Dataset — UCI Heart Disease (combined)

- 920 patients × 16 clinical columns.
- Pooled from four hospitals (multi-site).
- Target 'num' (0–4) → binarized: 0 = no disease, 1–4 = disease.
- Contains real missing values — handled honestly, not dropped.

304
Cleveland records

293
Hungary records

200
VA Long Beach records

123
Switzerland records

Outcome balance

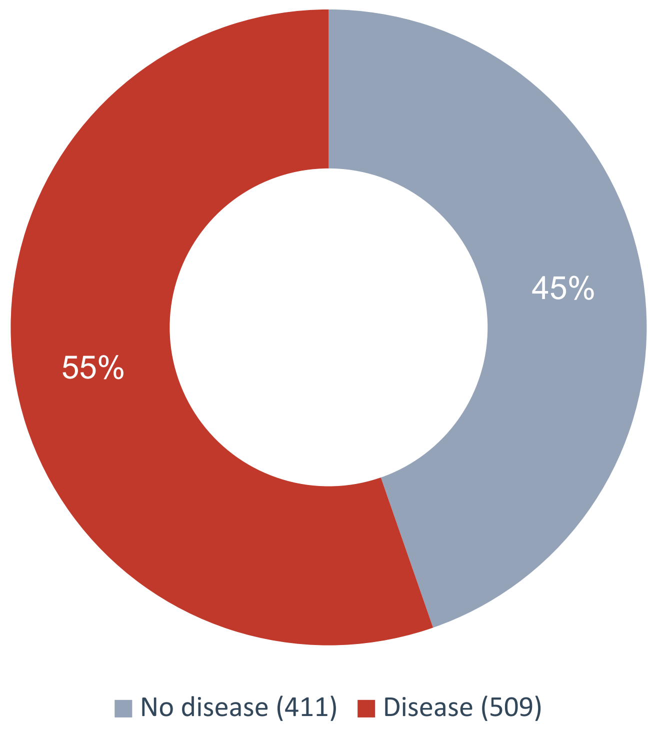

ROADMAP

Eight Steps, One Pipeline

A1 Inspect
shape, types, missingness

A2 Preprocess
clean · encode · scale · split

A3 Explore
EDA: what separates sick

A4 Baseline
Logistic Regression

A5 Compare
RF & Gradient Boosting

A6 Tune
bias-variance · reg · threshold

A7 Cluster
K-means risk phenotypes

A8 Save
persist + bridge to AI

STEP A1

Load & Inspect the Data

- Read the CSV into a pandas DataFrame — change nothing yet.
- Confirmed 920 rows × 16 columns; identified each column's meaning & type.
- Target distribution: 411 healthy vs 509 diseased (after binarizing).
- Surfaced the missing-value map (right) — the central challenge.

Key finding

ca (66%) and thal (53%) are missing for over half of patients — too valuable to drop, too sparse to ignore. This drove the A2 strategy.

Missing values by column (%)

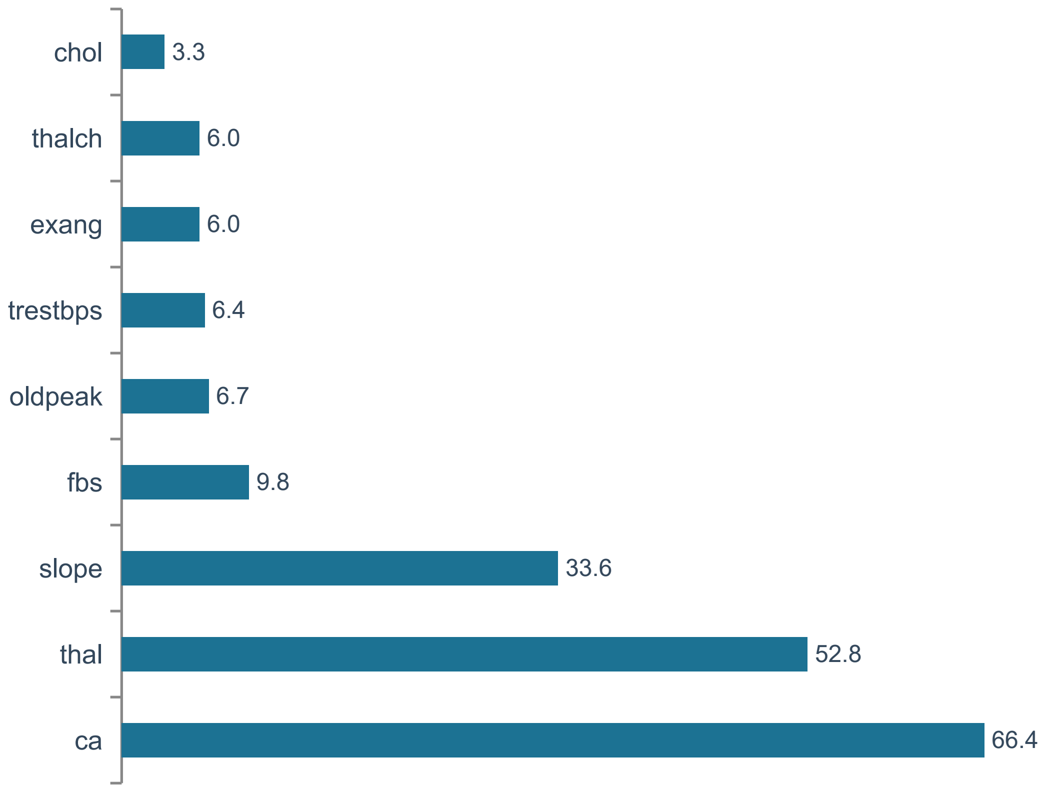

STEP A2

Preprocessing — Leakage-Safe Pipeline

- Binarize target: num 1–4 → 1 (disease present).
- Drop id (row number) and dataset (hospital — a confound).
- Fix hidden missingness: chol = 0 is impossible → 172 cases set to missing.
- Add ca_missing / thal_missing flags so the model knows a value was estimated.
- Median/mode impute · one-hot encode · standardize — all inside one Pipeline.
- Fit on TRAIN only, apply to TEST → no information leaks across the split.

736 / 184
stratified train / test

0.553 / 0.554
disease rate · train vs test

15 → 27
features after encoding

0
missing values remaining

Why leakage matters

If fill-values or scales were learned from the whole dataset, the test score would be inflated and dishonest.

STEP A3

Exploratory Analysis — What Separates Sick from Healthy

Mean value by outcome (training set)

Feature	Disease	Healthy	Signal
thalch (max HR)	128.2	149.1	strong
oldpeak	1.3	0.4	strong
ca (vessels)	1.1	0.2	strong
age	56.2	50.4	moderate
chol	254	240	weak
trestbps	134	130	weak

Silent danger

Asymptomatic chest pain shows the HIGHEST disease rate (78%) — the 'silent' group screening can miss.

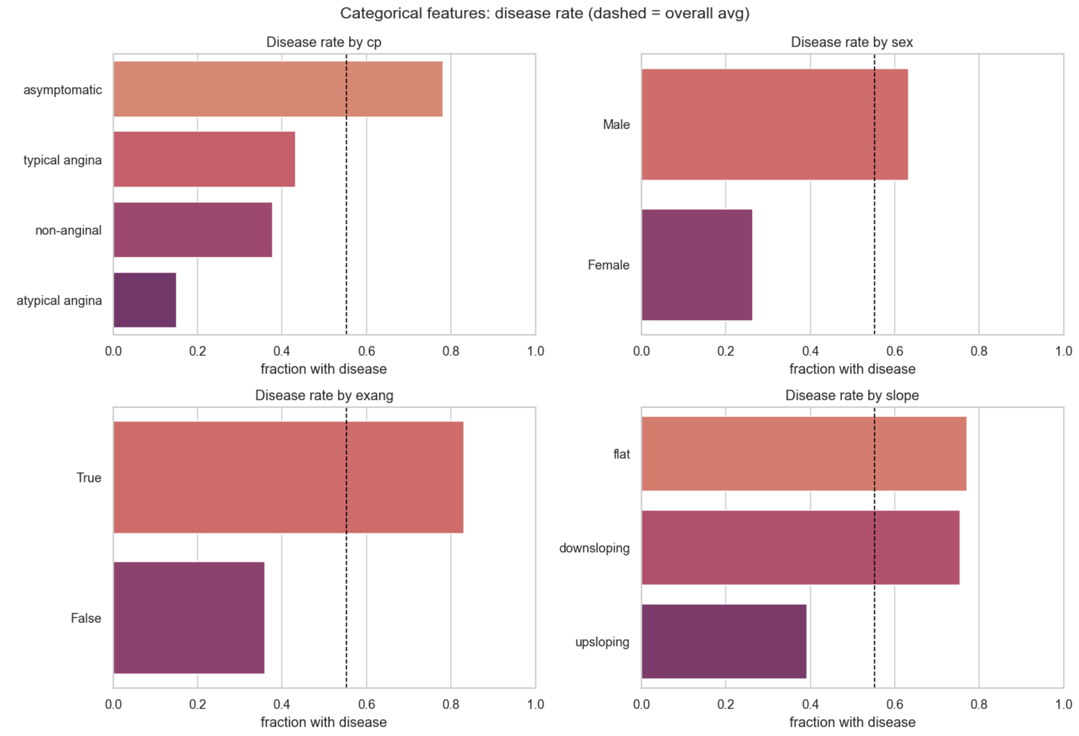

STEP A3

Feature Correlation with Disease

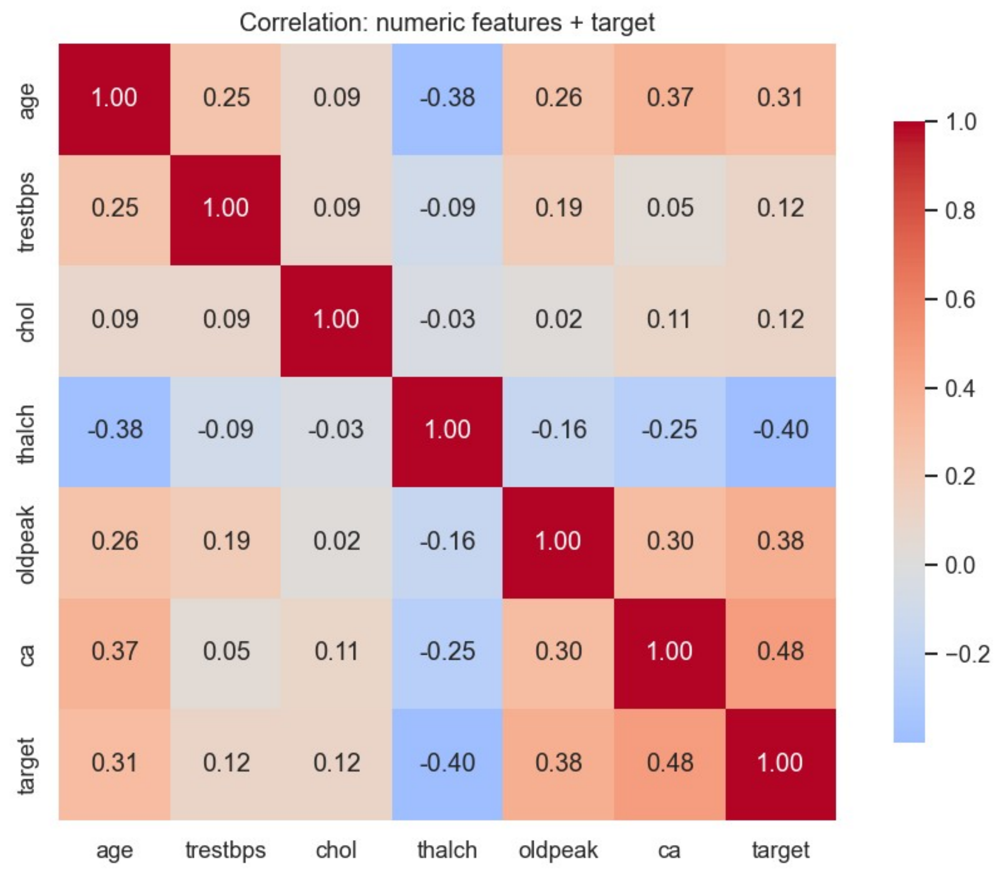

Strength of link with disease

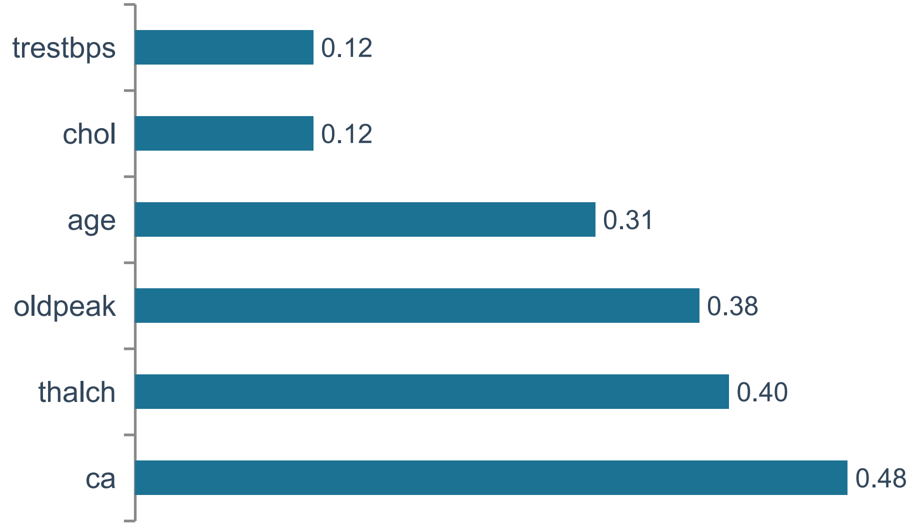

- ca, thalch, oldpeak, age are the strongest signals.
- chol & BP correlate weakly — famous ≠ predictive.
- No pair exceeds 0.38 → low multicollinearity (good for the baseline).

STEP A4

Baseline Model — Logistic Regression

0.799	Accuracy
0.804	Precision
0.843	Recall
0.744	Specificity
0.823	F1
0.897	ROC-AUC

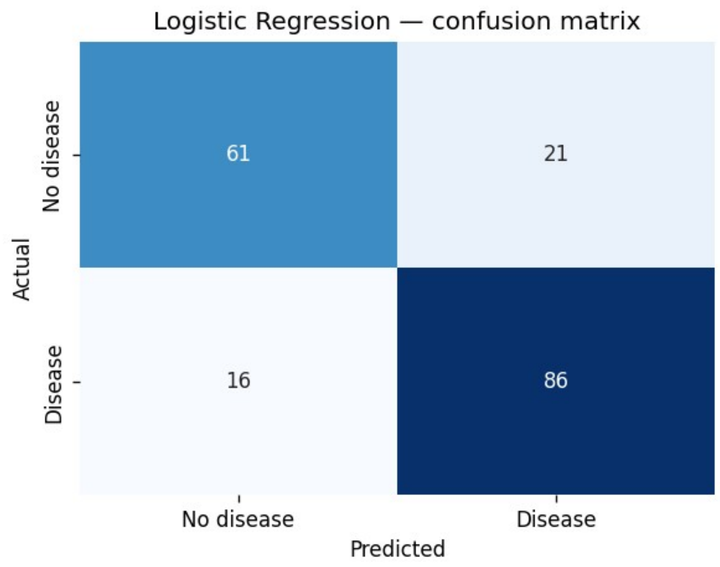

Read it clinically

Caught 86 of 102 sick patients; 16 missed cases (false negatives) are the target for tuning. Train 0.836 vs test 0.799 → small gap, a healthy baseline.

STEP A5

Three Models Compared

Model	Acc	Recall	Spec	F1	AUC	Train
Logistic Regression	0.799	0.843	0.744	0.823	0.897	0.836
Random Forest	0.810	0.863	0.744	0.834	0.907	1.000
Gradient Boosting ✓	0.821	**0.892**	0.732	**0.847**	0.905	0.916

- Both ensembles beat the baseline — real non-linear interactions exist.
- Random Forest hit a perfect 1.000 on training → clear overfitting.
- Winner: Gradient Boosting — best recall (0.892), F1 and accuracy; AUC tied for top.

ROC curves

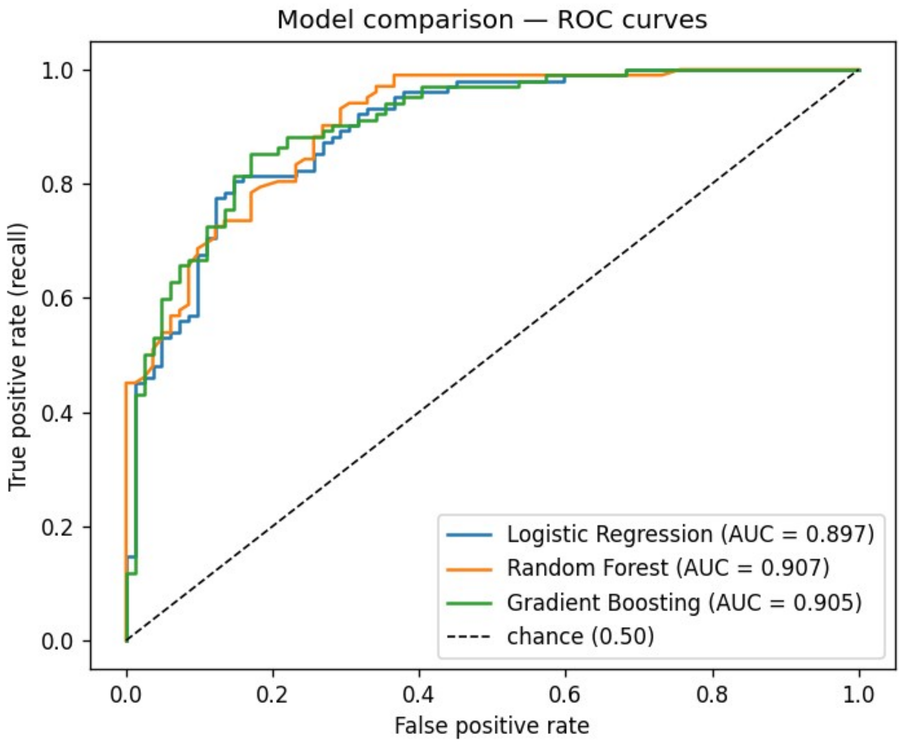

Bias–Variance Diagnosis (Learning Curves)

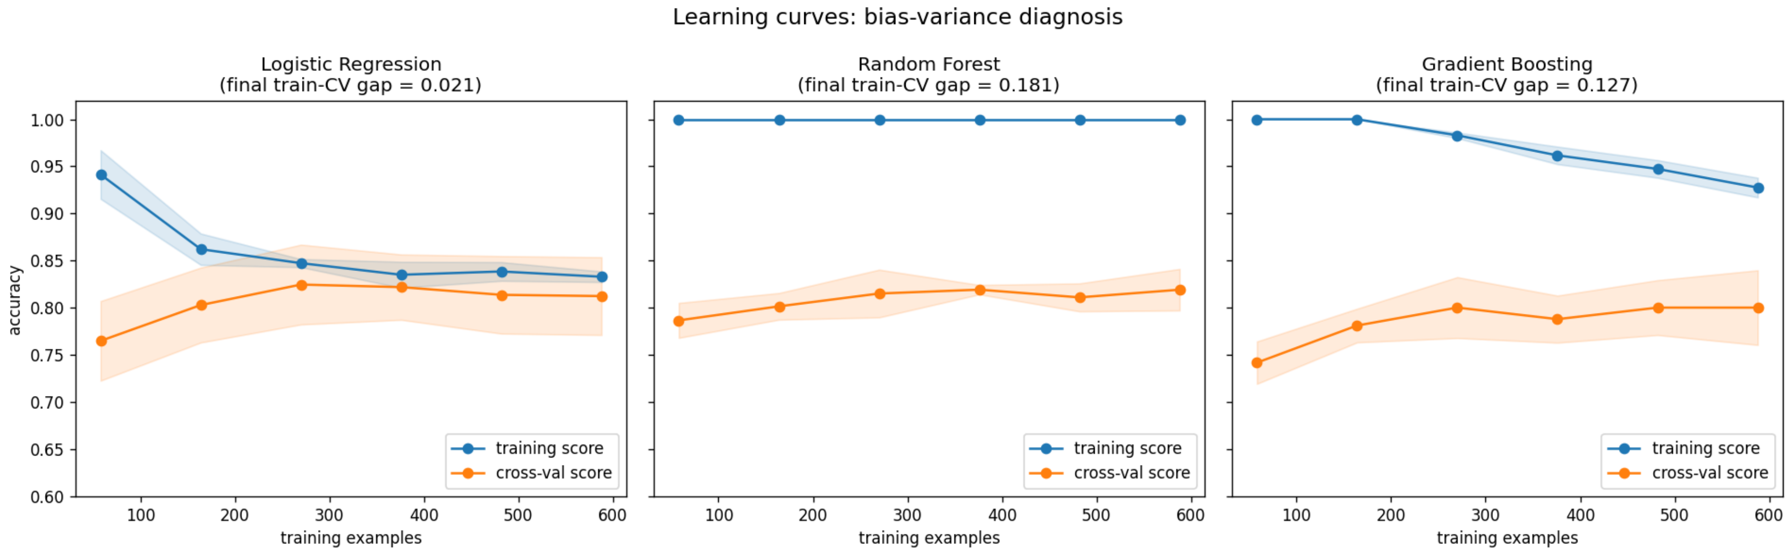

Logistic Regression

`gap 0.021`

Mild bias, low variance — curves meet; more data won't help.

Random Forest

`gap 0.181`

Severe overfitting — train pinned at 1.000 while CV lags.

Gradient Boosting

`gap 0.127`

Moderate, still converging — regularization should help.

STEP A6b

Regularization — Shrinking the Variance Gap

Train–CV accuracy gap: before vs after

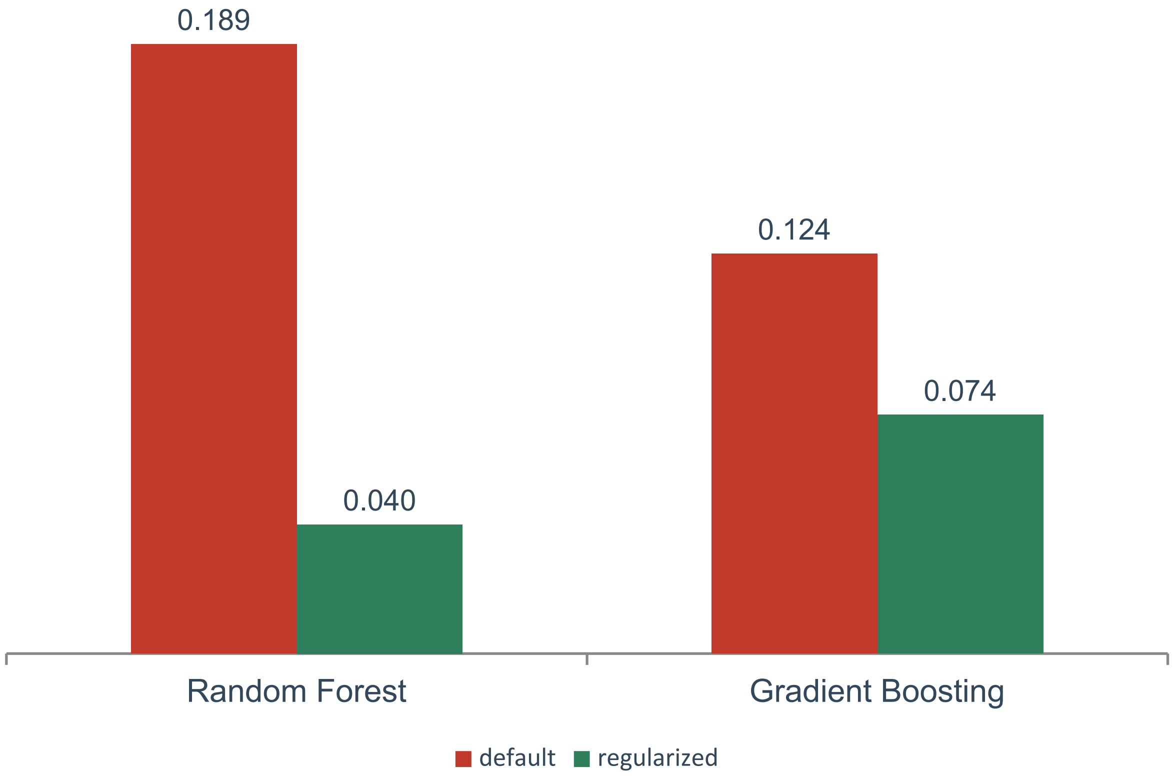

- Random Forest: gap crushed 0.189 → 0.040 (depth & leaf limits).
- Gradient Boosting: gap halved 0.124 → 0.074 — and CV accuracy rose.
- Logistic Regression is bias-limited: regularization couldn't help it.
- L1 (lasso) drove 11 of 27 coefficients to exactly zero — automatic feature selection.

The lesson

The cure depends on the diagnosis: high variance → constrain complexity; high bias → add capacity.

STEP A6c

Threshold Tuning — Prioritizing Recall

Default 0.50 vs tuned 0.35

Metric	0.50	0.35
Missed cases (FN)	10	**4**
Recall	0.902	**0.961**
Precision	0.821	0.778
Specificity	0.756	0.659

−60%

missed cases

Why lower?

A missed heart disease ≫ a false alarm. We accept 8 extra alarms to catch 6 more real cases.

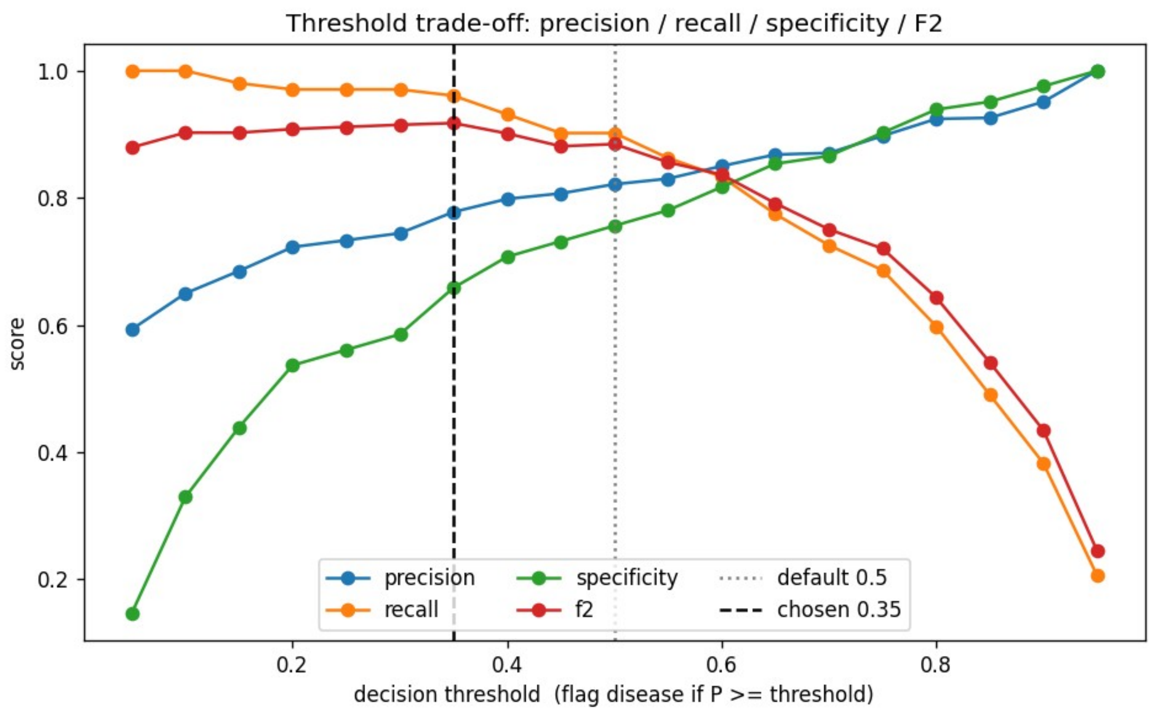

STEP A7

Risk Phenotypes — K-means (k = 4)

Pri	n	Disease	Phenotype
1	74	0.84	Advanced multi-vessel disease
2	211	0.74	Hypertensive-hypercholesterolemic
3	292	0.66	Reduced-exertion / low max-HR
4	343	0.28	Younger lower-risk

- Unsupervised: target hidden during fitting, used only to interpret.
- Chose k=4 for clinical granularity (silhouette was flat) — a transparent judgment call.
- Each phenotype gives the AI agent a priority hint beyond raw probability.

Clusters in 2D (PCA view)

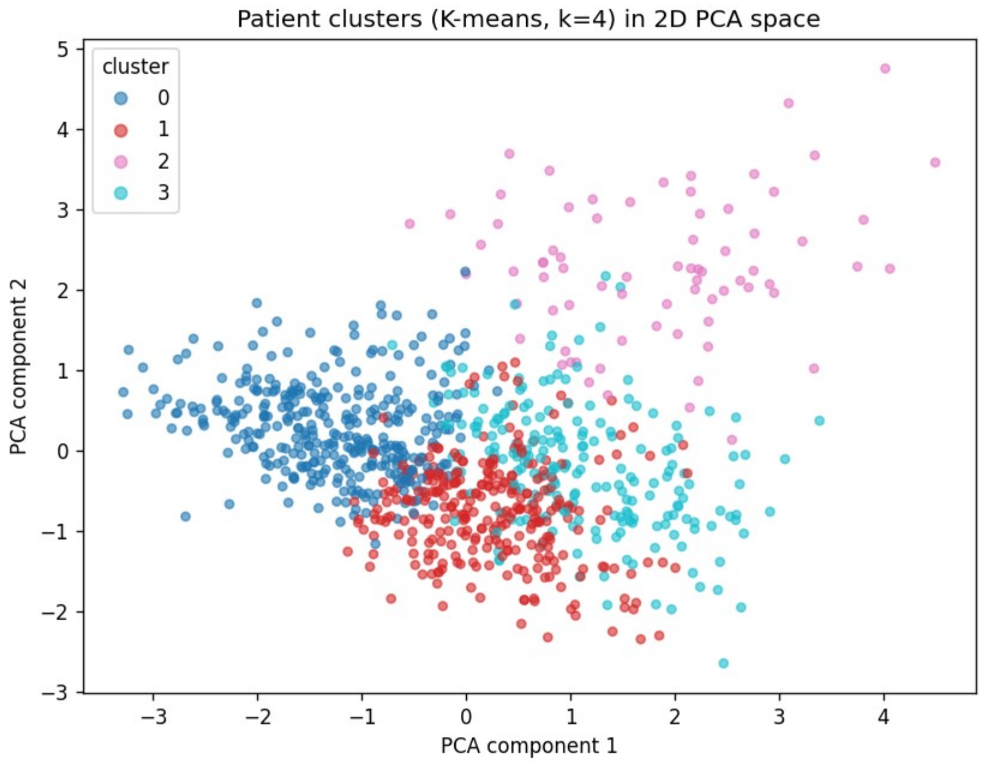

STEP A8

Persist & Bridge to the AI Phase

Saved artifacts (models/)

- risk_model.joblib — regularized Gradient Boosting pipeline (refit on all 920).
- phenotype_model.joblib — K-means impute+scale+cluster.
- metadata.json — threshold 0.35, feature list, phenotype map, test metrics.

0.961	0.908	0.826
Recall	ROC-AUC	Accuracy

Held-out test metrics @ threshold 0.35

```
model_api.assess(patient)

HIGH-risk patient
  probability : 0.976  → high-risk
  phenotype   : Advanced multi-vessel (priority 1)

LOW-risk patient
  probability : 0.013  → cleared
  phenotype   : Younger lower-risk (priority 4)
```

One clean call returns risk + phenotype — the launchpad for Phase B.

RESULTS

Phase A — Headline Results

0.908	0.961	4 / 102	4
ROC-AUC	Recall (tuned)	Missed cases	Risk phenotypes

Recall across models

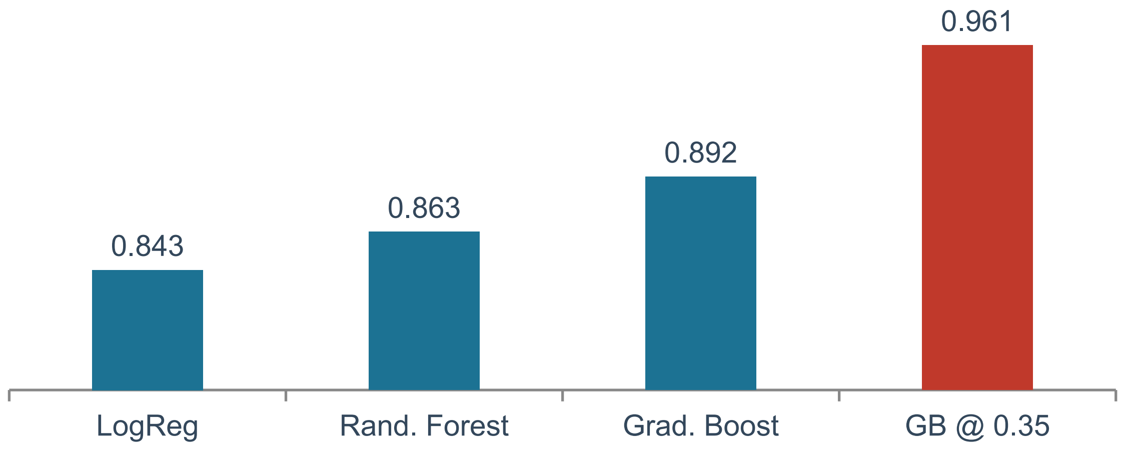

What we delivered

- Best model: regularized Gradient Boosting.
- Full leakage-safe pipeline, reused across every step.
- Tuned for safety: only 4 missed cases out of 102.
- Four interpretable phenotypes to guide triage.
- One-call interface bridging ML → AI.

RIGOR

Beyond the Baseline — Did We Leave Accuracy on the Table?

Three independent attempts to beat Gradient Boosting, all cross-validated

Approach	CV AUC	Test AUC	Test Acc	Test Recall	Outcome
Gradient Boosting (chosen) ✓	0.875	0.908	0.826	**0.961**	baseline
+ Hyperparameter search (+ KNN)	0.882	0.910	0.815	0.961	AUC ↑ = noise
+ Feature engineering	0.876	0.913	**0.837**	0.961	best test acc, CV flat
+ Neural network (MLP)	0.885	0.898	0.832	0.941	recall ↓, no gain

One ceiling, confirmed three ways

Every approach lands in the same 0.82–0.84 accuracy / ~0.90 AUC band. Where a single test split looked better (feature engineering, 0.837), 5-fold cross-validation did NOT confirm it — the tell-tale sign of noise, not signal. We kept Gradient Boosting for its best held-out recall and AUC, the metrics that matter for triage.

EDGE

Why This Project Has a Clear Edge

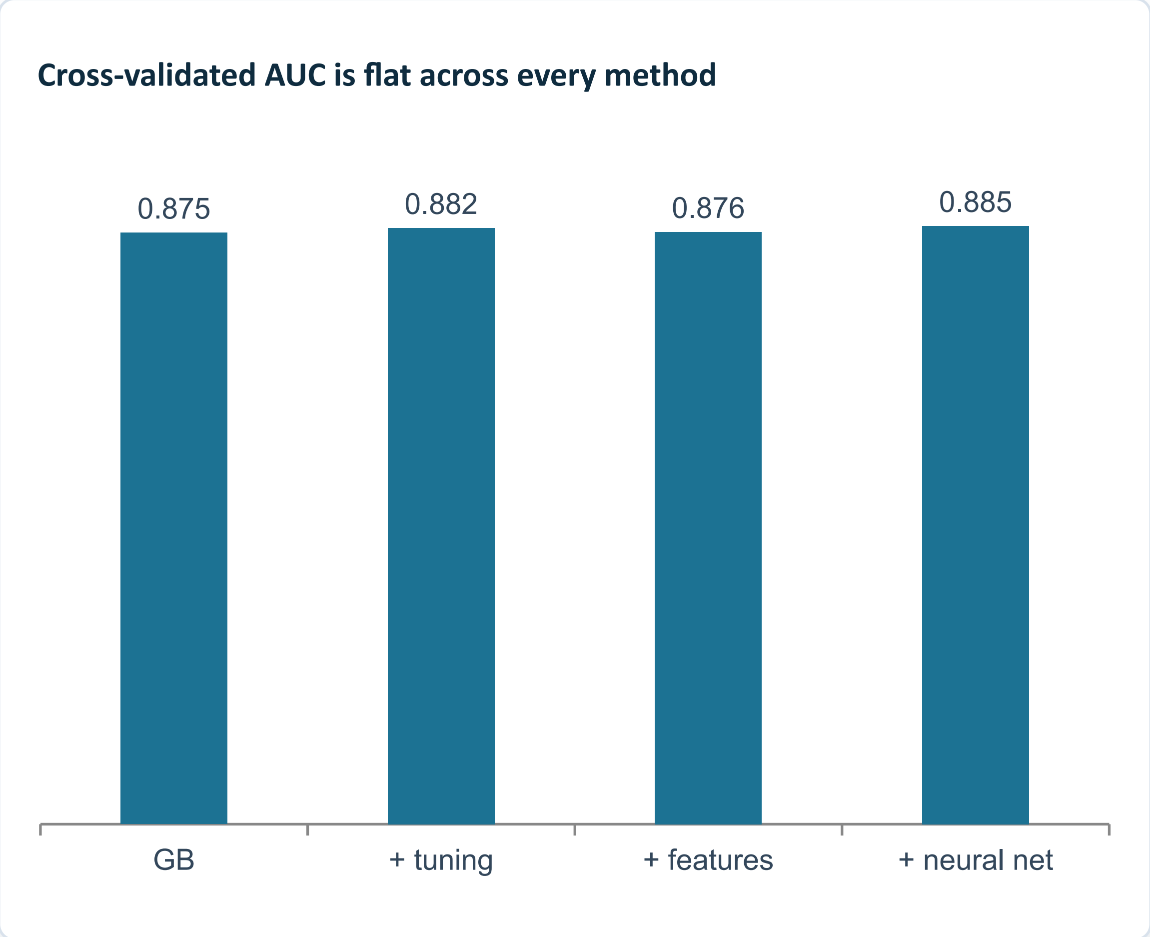

What sets this work apart

- Benchmarked 6 approaches — 3 model families + tuning + engineered features + a neural net. Most projects stop at one.
- Used cross-validation to reject test-set noise — refused to chase a fragile 0.837 accuracy.
- Proved the data's ceiling with evidence, not assumption.
- Optimized recall (clinical safety) over vanity accuracy.
- Leakage-safe pipeline end-to-end; fully reproducible.

Examiners reward understanding WHY, not just HOW — this project shows both.

Machine Learning Phase: Complete

A trained, evaluated, deployable cardiac-risk engine.

- Inspected, cleaned, and encoded 920 records — leakage-safe.
- Compared three models; chose Gradient Boosting on recall & AUC.
- Diagnosed bias/variance, regularized, and tuned the threshold for safety.
- Discovered four clinical risk phenotypes with K-means.
- Saved the model and exposed assess(patient) for the AI phase.

Next

Say “start the AI phase” to build A* search, the CSP allocator, the knowledge base, and the Gemini + LangChain agent — all on top of assess().